

# Assignment 4: Neural Networks

Due Date: November 17, 11:59 pm

A brief discussion of the assignment will take place in class on November 6. An online training session for PyTorch will be offered by TA Hazem Taha.

## Overview

In this assignment, you are required to train a neural network to solve a multi-class classification problem using both your own neural network implementation and using **PyTorch**. You will use the **Fashion MNIST** dataset, which contains  $28 \times 28$  grayscale images of articles of clothing and fashion accessories. Each image is assigned a label from ten possible classes.

## Requirements

### 1. Dataset

#### **Fashion MNIST Dataset:**

There are two ways to download the dataset:

1. **Through Kaggle:** The dataset can be downloaded from <https://www.kaggle.com/datasets/zalando-research/fashionmnist?resource=download>. The dataset will be in CSV files with the following content:
  - `fashion-mnist_train.csv`: 60,000 training examples.
  - `fashion-mnist_test.csv`: 10,000 test examples.

Each example is a  $28 \times 28$  grayscale image, resulting in 784 pixels per image. Each pixel-value is an integer between 0 and 255, indicating the lightness or darkness of that pixel. The CSV files have 785 columns: the first column is the class label, and the remaining 784 columns are the pixel-values.

2. **Using `datasets.FashionMNIST`:** The data will be downloaded as tensors using the following code:

```
1 from torchvision import datasets
2
3 train_dataset = datasets.FashionMNIST(root='./data', train=
  True, download=True)
4 test_dataset = datasets.FashionMNIST(root='./data', train=
  False, download=True)
```

### Data Split:

- The dataset is already split into a **training set of 60,000 examples** and a **test set of 10,000 examples**.
- You will further split the training data into a **training set** and a **validation set** (e.g., an 80/20 split).
- The validation set is needed for **early stopping** and for choosing your own hyperparameters.

## 2. Neural Network Implementations

You are required to implement **two neural networks**:

### 1. Custom Implementation:

- **Architecture:**
  - Input Layer: 784 units (raw pixel values).
  - **Hidden Layers:** At least **two hidden layers**.
    - \* Activation Function: Your choice (**ReLU is recommended**).
  - **Output Layer:** 10 units.
    - \* Activation Function: **Softmax**.
- **Training:**
  - Loss Function: **Cross-entropy loss**.
  - Techniques:
    - \* **Mini-batch gradient descent**.
    - \* **Weight decay** (L2 regularization).
    - \* **Early stopping**.

### 2. PyTorch Implementation:

- Replicate the **same neural network** architecture using **PyTorch**.
- Utilize **PyTorch functionalities** for defining the model, loss function, optimizer, and training loop.

## 3. Model Training and Evaluation

- Train **two models** with **different initializations** using each of the **two implementations** discussed above (total of **four models**).
- Plot the **learning curves** (**cross-entropy loss vs. number of epochs**) for both training and validation sets.
- Record the **training and test misclassification errors** for each model.
- Organize the results in a clear table.
- Specify all **hyperparameter** values used and the activation function used at the hidden units.
- Include any other observations you find valuable.

# Guidelines

## 1. Code Development

- **Modularity:** Write functions for each of the main tasks and for any repeated operations.
- **Vectorization:** Use vectorized operations with NumPy wherever possible for efficiency.
- **Comments and Documentation:** Include instructive comments explaining your code.
- **Libraries:**
  - **Allowed:**
    - \* **NumPy** for linear algebra and array manipulation.
    - \* **scikit-learn** or **PyTorch** (for data manipulation tasks only, not for **neural network implementation** in your custom code) (e.g., **shuffling, splitting, standardization**).
    - \* **Matplotlib** or similar libraries for plotting.
  - **Not Allowed in Custom Implementation:**
    - \* Using **PyTorch** or any high-level neural network libraries (e.g., Keras) for implementing the neural network in your custom implementation.

## 2. Randomization

- Whenever you use **randomization** in your code, use a number **formed with the last 4 digits of your student ID**, in any order, as the **seed** for the pseudo number generator.
- Ensure that the seed is set for **all random number generators** used (e.g., NumPy, PyTorch).

## 3. Report Writing

Your report should include:

- **Introduction:** Brief overview of the assignment objectives.
- **Methodology:** Description of neural network architectures, training processes, and hyperparameters.
- **Results:**
  - Learning curves for each model (training and validation loss over epochs).
  - Training and test misclassification errors.
  - A table summarizing the results for all models.

- **Discussion:** Analysis of results, impact of different initializations and hyperparameters, comparison between custom and PyTorch implementations.
- **Conclusion:** Summary of findings and suggestions for potential improvements if more time was available.

## 4. Demo Video

- **Duration:** Maximum of 1 minute.
- **Content:** Scroll through your code and briefly explain what each part does.
- **Format:** Ensure the video is clear and audible. Acceptable formats include MP4, AVI, or MOV.

## Deliverables

Submit the following by **November 17, 11:59 pm**:

1. **Report:** A PDF file containing your detailed report.
2. **Code Files:** Python files (.py) of your implementations.
  - Your code should include the implementations of the four neural networks as discussed (two custom implementations with different initializations and two PyTorch implementations).
  - Include any necessary instructions to run your code.
3. **Demo Video:** A video file showcasing your code with explanations.

## Additional Notes

### 1. Hyperparameters

You may use the following suggested hyperparameters or search for better ones:

Hidden Layers	Hidden Units	Learning Rate	Batch Size	Weight Decay ( $\lambda$ )
2	156	0.01	128	0.0018738
3	92	0.01	128	0.00231
4	80	0.01	128	0.001232

1

<sup>1</sup>These values are based on Hazem Taha's implementation. The test misclassification rate obtained with the models trained with these values was below 14%.

## 2. Data Preprocessing

Below is a sample code reference for loading and preprocessing the data using `datasets.FashionMNIST`:

```
1 import numpy as np
2 import torch
3 from torchvision import datasets
4
5 # Seed the pseudo number generators (e.g., if your ID ends with
   0393)
6 np.random.seed(3093)
7 torch.manual_seed(3093)
8
9 # Define the number of classes
10 num_classes = 10 # For Fashion MNIST
11
12 # Load Fashion MNIST dataset
13 train_dataset = datasets.FashionMNIST(root='./data', train=True,
   download=True)
14 test_dataset = datasets.FashionMNIST(root='./data', train=False,
   download=True)
15
16 print(train_dataset.data.shape, test_dataset.data.shape)
17
18 # Prepare the data as numpy arrays
19 X_train = train_dataset.data.numpy().reshape(-1, 28 * 28).astype(
   'float32') / 255.0
20 Y_train = train_dataset.targets.numpy()
21
22 X_test = test_dataset.data.numpy().reshape(-1, 28 * 28).astype('
   float32') / 255.0
23 Y_test = test_dataset.targets.numpy()
24
25 # Split the training set into train and validation sets (80% /
   20%)
26 validation_size = int(0.2 * X_train.shape[0])
27 X_validation, Y_validation = X_train[:validation_size], Y_train[:
   validation_size]
28 X_train, Y_train = X_train[validation_size:], Y_train[
   validation_size:]
29
30 # Save original labels before one-hot encoding
31 Y_train_orig = Y_train
32 Y_validation_orig = Y_validation
33 Y_test_orig = Y_test
34
35 # Convert labels to one-hot encoding for multi-class
   classification
36 def one_hot_encode(labels, num_classes):
37     return np.eye(num_classes)[labels]
38
```

```
39 Y_train = one_hot_encode(Y_train, num_classes)
40 Y_validation = one_hot_encode(Y_validation, num_classes)
41 Y_test = one_hot_encode(Y_test, num_classes)
42
43 # Standardizing the data
44
45 # Calculate the mean and standard deviation of the training
    features
46 X_train_mean = X_train.mean(axis=0)
47 X_train_std = X_train.std(axis=0)
48 X_train_std[X_train_std == 0] = 1 # To avoid division by zero
49
50 # Standardize all three subsets of data
51 X_train = (X_train - X_train_mean) / X_train_std
52 X_validation = (X_validation - X_train_mean) / X_train_std
53 X_test = (X_test - X_train_mean) / X_train_std
```

### 3. Performance Benchmark

- Models trained with the suggested hyperparameters achieved a test misclassification rate below 14%.

### 4. Improvement Suggestions

- You may experiment with different hyperparameters, activation functions, or architectures to improve performance.
- Document any changes and their impact on the results.

### 5. Academic Integrity

- Ensure all submitted work is your own.
- Cite any external resources or code snippets used.
- Adhere to the university's policies on academic integrity.